

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt,seaborn as sns
```

```
In [2]: train_df=pd.read_csv(r"C:\Users\yasoda\Documents\202U1A05C1\Mobile_Price_Classification_test.csv")
train_df
```

Out[2]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w
0	1	1043	1	1.8	1	14	0	5	0.1	193	...	16	226	1412	3476	12	
1	2	841	1	0.5	1	4	1	61	0.8	191	...	12	746	857	3895	6	
2	3	1807	1	2.8	0	1	0	27	0.9	186	...	4	1270	1366	2396	17	
3	4	1546	0	0.5	1	18	1	25	0.5	96	...	20	295	1752	3893	10	
4	5	1434	0	1.4	0	11	1	49	0.5	108	...	18	749	810	1773	15	
...
995	996	1700	1	1.9	0	0	1	54	0.5	170	...	17	644	913	2121	14	
996	997	609	0	1.8	1	0	0	13	0.9	186	...	2	1152	1632	1933	8	
997	998	1185	0	1.4	0	1	1	8	0.5	80	...	12	477	825	1223	5	
998	999	1533	1	0.5	1	0	0	50	0.4	171	...	12	38	832	2509	15	
999	1000	1270	1	0.5	0	4	1	35	0.1	140	...	19	457	608	2828	9	

1000 rows × 21 columns



```
In [3]: test_df=pd.read_csv(r"C:\Users\yasoda\Documents\202U1A05C1\Mobile_Price_Classification_train.csv")
test_df
```

Out[3]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w
0	842	0	2.2	0	1	0	7	0.6	188	2	...	20	756	2549	9	
1	1021	1	0.5	1	0	1	53	0.7	136	3	...	905	1988	2631	17	
2	563	1	0.5	1	2	1	41	0.9	145	5	...	1263	1716	2603	11	
3	615	1	2.5	0	0	0	10	0.8	131	6	...	1216	1786	2769	16	
4	1821	1	1.2	0	13	1	44	0.6	141	2	...	1208	1212	1411	8	
...
1995	794	1	0.5	1	0	1	2	0.8	106	6	...	1222	1890	668	13	
1996	1965	1	2.6	1	0	0	39	0.2	187	4	...	915	1965	2032	11	
1997	1911	0	0.9	1	1	1	36	0.7	108	8	...	868	1632	3057	9	
1998	1512	0	0.9	0	4	1	46	0.1	145	5	...	336	670	869	18	
1999	510	1	2.0	1	5	1	45	0.9	168	6	...	483	754	3919	19	

2000 rows × 21 columns



```
In [4]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               1000 non-null   int64
1   battery_power    1000 non-null   int64
2   blue             1000 non-null   int64
3   clock_speed      1000 non-null   float64
4   dual_sim         1000 non-null   int64
5   fc               1000 non-null   int64
6   four_g           1000 non-null   int64
7   int_memory       1000 non-null   int64
8   m_dep            1000 non-null   float64
9   mobile_wt        1000 non-null   int64
10  n_cores          1000 non-null   int64
11  pc               1000 non-null   int64
12  px_height        1000 non-null   int64
13  px_width         1000 non-null   int64
14  ram              1000 non-null   int64
15  sc_h             1000 non-null   int64
16  sc_w             1000 non-null   int64
17  talk_time        1000 non-null   int64
18  three_g          1000 non-null   int64
19  touch_screen     1000 non-null   int64
20  wifi             1000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

```
In [5]: test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   battery_power   2000 non-null   int64  
 1   blue             2000 non-null   int64  
 2   clock_speed     2000 non-null   float64 
 3   dual_sim        2000 non-null   int64  
 4   fc              2000 non-null   int64  
 5   four_g          2000 non-null   int64  
 6   int_memory      2000 non-null   int64  
 7   m_dep           2000 non-null   float64 
 8   mobile_wt       2000 non-null   int64  
 9   n_cores         2000 non-null   int64  
10   pc              2000 non-null   int64  
11   px_height       2000 non-null   int64  
12   px_width        2000 non-null   int64  
13   ram             2000 non-null   int64  
14   sc_h            2000 non-null   int64  
15   sc_w            2000 non-null   int64  
16   talk_time       2000 non-null   int64  
17   three_g         2000 non-null   int64  
18   touch_screen    2000 non-null   int64  
19   wifi            2000 non-null   int64  
20   price_range     2000 non-null   int64  
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
```

```
In [6]: x=train_df.drop('wifi',axis=1)
        y=train_df['wifi']
```

```
In [7]: x=test_df.drop('wifi',axis=1)
        y=test_df['wifi']
```

```
In [8]: train_df['dual_sim'].value_counts()
```

```
Out[8]: dual_sim
1      517
0      483
Name: count, dtype: int64
```

```
In [9]: test_df['blue'].value_counts()
```

```
Out[9]: blue
0      1010
1       990
Name: count, dtype: int64
```

```
In [10]: T={"Home Owner":{"Yes":1,"No":0}}  
train_df=train_df.replace(T)  
print(train_df)
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	
0	1	1043	1	1.8	1	14	0	5	\
1	2	841	1	0.5	1	4	1	61	
2	3	1807	1	2.8	0	1	0	27	
3	4	1546	0	0.5	1	18	1	25	
4	5	1434	0	1.4	0	11	1	49	
..	
995	996	1700	1	1.9	0	0	1	54	
996	997	609	0	1.8	1	0	0	13	
997	998	1185	0	1.4	0	1	1	8	
998	999	1533	1	0.5	1	0	0	50	
999	1000	1270	1	0.5	0	4	1	35	

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w	
0	0.1	193	...	16	226	1412	3476	12	7	\
1	0.8	191	...	12	746	857	3895	6	0	
2	0.9	186	...	4	1270	1366	2396	17	10	
3	0.5	96	...	20	295	1752	3893	10	0	
4	0.5	108	...	18	749	810	1773	15	8	
..	
995	0.5	170	...	17	644	913	2121	14	8	
996	0.9	186	...	2	1152	1632	1933	8	1	
997	0.5	80	...	12	477	825	1223	5	0	
998	0.4	171	...	12	38	832	2509	15	11	
999	0.1	140	...	19	457	608	2828	9	2	

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

```
In [11]: T={"Home Owner":{"Yes":1,"No":0}}  
train_df=train_df.replace(T)  
print(train_df)
```


	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	1	1043	1	1.8	1	14	0	5
1	2	841	1	0.5	1	4	1	61
2	3	1807	1	2.8	0	1	0	27
3	4	1546	0	0.5	1	18	1	25
4	5	1434	0	1.4	0	11	1	49
..
995	996	1700	1	1.9	0	0	1	54
996	997	609	0	1.8	1	0	0	13
997	998	1185	0	1.4	0	1	1	8
998	999	1533	1	0.5	1	0	0	50
999	1000	1270	1	0.5	0	4	1	35

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w
0	0.1	193	...	16	226	1412	3476	12	7
1	0.8	191	...	12	746	857	3895	6	0
2	0.9	186	...	4	1270	1366	2396	17	10
3	0.5	96	...	20	295	1752	3893	10	0
4	0.5	108	...	18	749	810	1773	15	8
..
995	0.5	170	...	17	644	913	2121	14	8
996	0.9	186	...	2	1152	1632	1933	8	1
997	0.5	80	...	12	477	825	1223	5	0
998	0.4	171	...	12	38	832	2509	15	11
999	0.1	140	...	19	457	608	2828	9	2

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

```
In [12]: T={"Home Owner":{"Yes":1,"No":0}}  
test_df=test_df.replace(T)  
print(test_df)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	842	0	2.2	0	1	0	7
1	1021	1	0.5	1	0	1	53
2	563	1	0.5	1	2	1	41
3	615	1	2.5	0	0	0	10
4	1821	1	1.2	0	13	1	44
...
1995	794	1	0.5	1	0	1	2
1996	1965	1	2.6	1	0	0	39
1997	1911	0	0.9	1	1	1	36
1998	1512	0	0.9	0	4	1	46
1999	510	1	2.0	1	5	1	45

	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w
0	0.6	188	2	...	20	756	2549	9	7
1	0.7	136	3	...	905	1988	2631	17	3
2	0.9	145	5	...	1263	1716	2603	11	2
3	0.8	131	6	...	1216	1786	2769	16	8
4	0.6	141	2	...	1208	1212	1411	8	2
...
1995	0.8	106	6	...	1222	1890	668	13	4
1996	0.2	187	4	...	915	1965	2032	11	10
1997	0.7	108	8	...	868	1632	3057	9	1
1998	0.1	145	5	...	336	670	869	18	10
1999	0.9	168	6	...	483	754	3919	19	4

	talk_time	three_g	touch_screen	wifi	price_range
0	19	0	0	1	1
1	7	1	1	0	2
2	9	1	1	0	2
3	11	1	0	0	2
4	15	1	1	0	1
...
1995	19	1	1	0	0
1996	16	1	1	1	2
1997	5	1	1	0	3
1998	19	1	1	1	0
1999	2	1	1	1	3

[2000 rows x 21 columns]

```
In [13]: x=train_df.drop('wifi',axis=1)
y=train_df['wifi']
```

```
In [14]: x=test_df.drop('wifi',axis=1)
y=test_df['wifi']
```

```
In [15]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,train_size=0.7,random_state=42)
x_train.shape,x_test.shape
```

```
Out[15]: ((1400, 20), (600, 20))
```

```
In [16]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[16]: ▾ RandomForestClassifier
RandomForestClassifier()
```

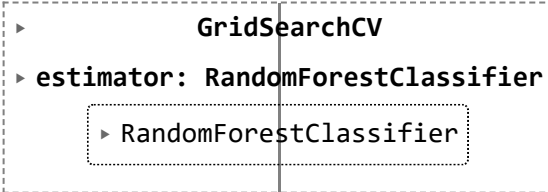
In a Jupyter Environment ,please rerun this cell to show the HTML representation

```
In [17]: rf = RandomForestClassifier()
```

```
In [18]: params = {'max_depth': [2,3,5,10,20],
'min_samples_leaf': [5,10,20,50,100,200],
'n_estimators': [10,25,30,50,100,200]}
```

```
In [19]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rf,param_grid=params,cv = 2, scoring='accuracy')
grid_search.fit(x_train,y_train)
```

```
Out[19]:
```



```
  ▸ GridSearchCV
  ▸ estimator: RandomForestClassifier
    ▸ RandomForestClassifier
```

In a Jupyter Environment, please rerun this cell to show the HTML representation

```
In [20]: grid_search.best_score_
```


```
Out[20]: 0.5228571428571429
```

```
In [21]: rf_best = grid_search.best_estimator_
print(rf_best)
```

```
RandomForestClassifier(max_depth=5, min_samples_leaf=100, n_estimators=50)
```

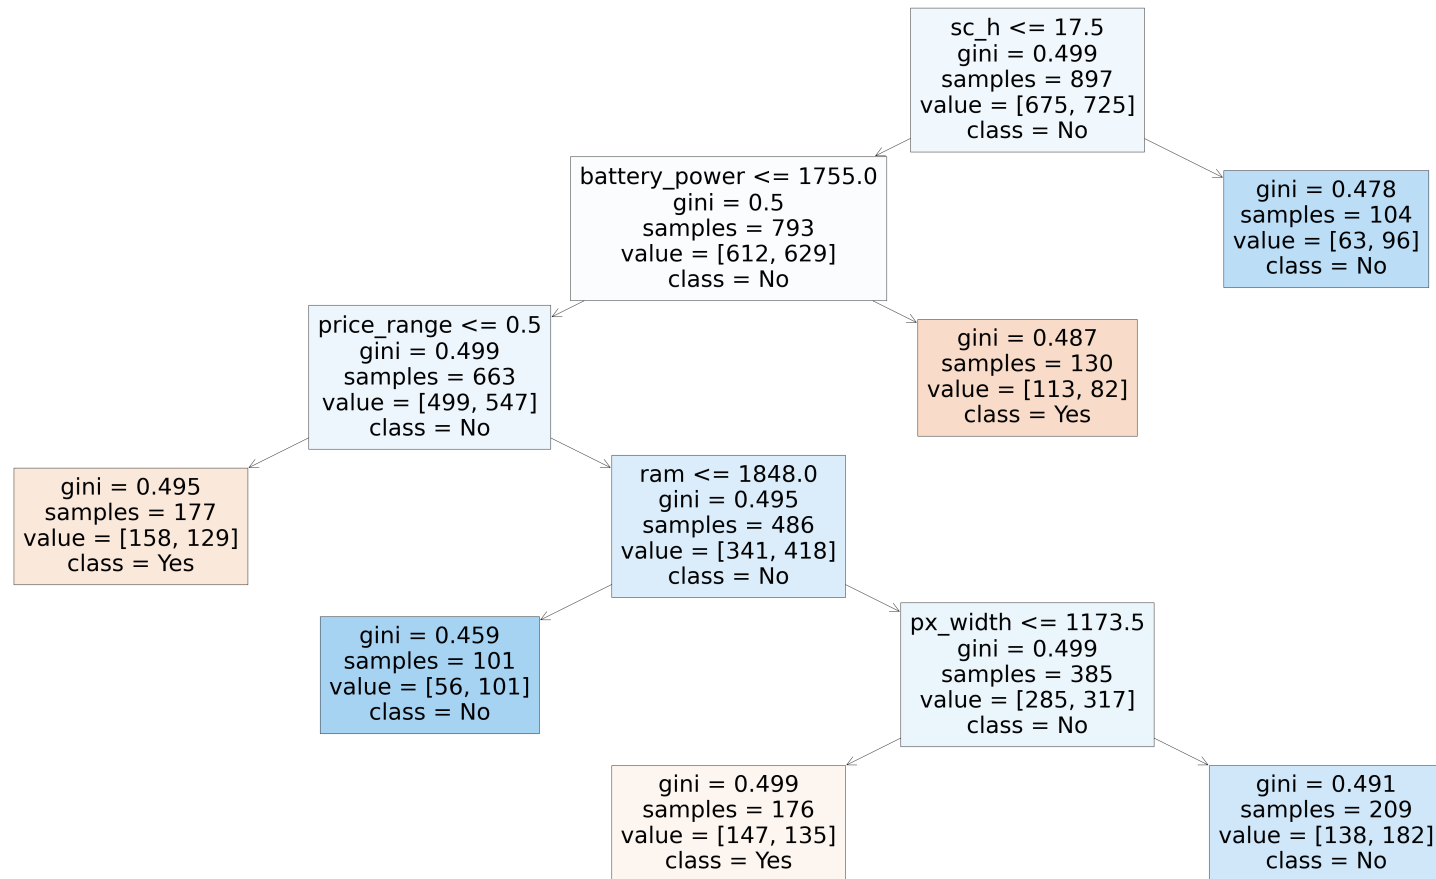
```
In [22]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[5], feature_names = x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[22]: [Text(0.5909090909090909, 0.875, 'ram <= 2731.5\ngini = 0.499\nsamples = 864\nvalue = [723, 677]\nclass = Yes'),
Text(0.36363636363636365, 0.625, 'talk_time <= 10.5\ngini = 0.495\nsamples = 571\nvalue = [501, 408]\nclass = Yes'),
Text(0.18181818181818182, 0.375, 'ram <= 1848.0\ngini = 0.5\nsamples = 269\nvalue = [212, 209]\nclass = Yes'),
Text(0.09090909090909091, 0.125, 'gini = 0.499\nsamples = 169\nvalue = [124, 138]\nclass = No'),
Text(0.2727272727272727, 0.125, 'gini = 0.494\nsamples = 100\nvalue = [88, 71]\nclass = Yes'),
Text(0.5454545454545454, 0.375, 'talk_time <= 14.5\ngini = 0.483\nsamples = 302\nvalue = [289, 199]\nclass = Yes'),
Text(0.45454545454545453, 0.125, 'gini = 0.444\nsamples = 108\nvalue = [116, 58]\nclass = Yes'),
Text(0.6363636363636364, 0.125, 'gini = 0.495\nsamples = 194\nvalue = [173, 141]\nclass = Yes'),
Text(0.8181818181818182, 0.625, 'sc_w <= 5.5\ngini = 0.495\nsamples = 293\nvalue = [222, 269]\nclass = No'),
Text(0.7272727272727273, 0.375, 'gini = 0.5\nsamples = 166\nvalue = [140, 144]\nclass = No'),
Text(0.9090909090909091, 0.375, 'gini = 0.478\nsamples = 127\nvalue = [82, 125]\nclass = No')]
```



```
In [23]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=["Yes","No"],filled=True)
```

```
Out[23]: [Text(0.6666666666666666, 0.9166666666666666, 'sc_h <= 17.5\ngini = 0.499\nsamples = 897\nvalue = [675, 725]\nclass = No'),
Text(0.5, 0.75, 'battery_power <= 1755.0\ngini = 0.5\nsamples = 793\nvalue = [612, 629]\nclass = No'),
Text(0.3333333333333333, 0.5833333333333334, 'price_range <= 0.5\ngini = 0.499\nsamples = 663\nvalue = [499, 547]\nclass = No'),
Text(0.16666666666666666, 0.41666666666666667, 'gini = 0.495\nsamples = 177\nvalue = [158, 129]\nclass = Yes'),
Text(0.5, 0.41666666666666667, 'ram <= 1848.0\ngini = 0.495\nsamples = 486\nvalue = [341, 418]\nclass = No'),
Text(0.3333333333333333, 0.25, 'gini = 0.459\nsamples = 101\nvalue = [56, 101]\nclass = No'),
Text(0.6666666666666666, 0.25, 'px_width <= 1173.5\ngini = 0.499\nsamples = 385\nvalue = [285, 317]\nclass = No'),
Text(0.5, 0.08333333333333333, 'gini = 0.499\nsamples = 176\nvalue = [147, 135]\nclass = Yes'),
Text(0.8333333333333334, 0.08333333333333333, 'gini = 0.491\nsamples = 209\nvalue = [138, 182]\nclass = No'),
Text(0.6666666666666666, 0.5833333333333334, 'gini = 0.487\nsamples = 130\nvalue = [113, 82]\nclass = Yes'),
Text(0.8333333333333334, 0.75, 'gini = 0.478\nsamples = 104\nvalue = [63, 96]\nclass = No')]
```

In [24]: rf_best.feature_importances_

Out[24]: array([0.07717307, 0.00879 , 0.04434522, 0.01953574, 0.03032118,
0.01928442, 0.08756478, 0.03099338, 0.05000981, 0.02300338,
0.06283458, 0.16018601, 0.07803434, 0.06519901, 0.05008368,
0.04774743, 0.10287966, 0.00509412, 0.00930888, 0.02761129])

```
In [25]: imp_df = pd.DataFrame({"Vername": x_train.columns, "Imp": rf_best.feature_importances_})  
imp_df.sort_values(by="Imp", ascending=False)
```

Out[25]:

	Vername	Imp
11	px_height	0.160186
16	talk_time	0.102880
6	int_memory	0.087565
12	px_width	0.078034
0	battery_power	0.077173
13	ram	0.065199
10	pc	0.062835
14	sc_h	0.050084
8	mobile_wt	0.050010
15	sc_w	0.047747
2	clock_speed	0.044345
7	m_dep	0.030993
4	fc	0.030321
19	price_range	0.027611
9	n_cores	0.023003
3	dual_sim	0.019536
5	four_g	0.019284
18	touch_screen	0.009309
1	blue	0.008790
17	three_g	0.005094

