

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt,seaborn as sns
```

```
In [2]: train_df=pd.read_csv(r"C:\Users\yasoda\Documents\202U1A05C1\Mobile_Price_Classification_test.csv")
train_df
```

```
Out[2]:
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w
0	1	1043	1	1.8	1	14	0	5	0.1	193	...	16	226	1412	3476	12	
1	2	841	1	0.5	1	4	1	61	0.8	191	...	12	746	857	3895	6	
2	3	1807	1	2.8	0	1	0	27	0.9	186	...	4	1270	1366	2396	17	
3	4	1546	0	0.5	1	18	1	25	0.5	96	...	20	295	1752	3893	10	
4	5	1434	0	1.4	0	11	1	49	0.5	108	...	18	749	810	1773	15	
...
995	996	1700	1	1.9	0	0	1	54	0.5	170	...	17	644	913	2121	14	
996	997	609	0	1.8	1	0	0	13	0.9	186	...	2	1152	1632	1933	8	
997	998	1185	0	1.4	0	1	1	8	0.5	80	...	12	477	825	1223	5	
998	999	1533	1	0.5	1	0	0	50	0.4	171	...	12	38	832	2509	15	
999	1000	1270	1	0.5	0	4	1	35	0.1	140	...	19	457	608	2828	9	

1000 rows × 21 columns



```
In [3]: test_df=pd.read_csv(r"C:\Users\yasoda\Documents\202U1A05C1\Mobile_Price_Classification_train.csv")
test_df
```

```
Out[3]:
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w
0	842	0	2.2	0	1	0	7	0.6	188	2	...	20	756	2549	9	
1	1021	1	0.5	1	0	1	53	0.7	136	3	...	905	1988	2631	17	
2	563	1	0.5	1	2	1	41	0.9	145	5	...	1263	1716	2603	11	
3	615	1	2.5	0	0	0	10	0.8	131	6	...	1216	1786	2769	16	
4	1821	1	1.2	0	13	1	44	0.6	141	2	...	1208	1212	1411	8	
...
1995	794	1	0.5	1	0	1	2	0.8	106	6	...	1222	1890	668	13	
1996	1965	1	2.6	1	0	0	39	0.2	187	4	...	915	1965	2032	11	
1997	1911	0	0.9	1	1	1	36	0.7	108	8	...	868	1632	3057	9	
1998	1512	0	0.9	0	4	1	46	0.1	145	5	...	336	670	869	18	
1999	510	1	2.0	1	5	1	45	0.9	168	6	...	483	754	3919	19	

2000 rows × 21 columns



```
In [4]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     1000 non-null   int64
1   battery_power          1000 non-null   int64
2   blue                   1000 non-null   int64
3   clock_speed            1000 non-null   float64
4   dual_sim               1000 non-null   int64
5   fc                     1000 non-null   int64
6   four_g                 1000 non-null   int64
7   int_memory             1000 non-null   int64
8   m_dep                  1000 non-null   float64
9   mobile_wt              1000 non-null   int64
10  n_cores                 1000 non-null   int64
11  pc                      1000 non-null   int64
12  px_height              1000 non-null   int64
13  px_width               1000 non-null   int64
14  ram                     1000 non-null   int64
15  sc_h                   1000 non-null   int64
16  sc_w                   1000 non-null   int64
17  talk_time              1000 non-null   int64
18  three_g                1000 non-null   int64
19  touch_screen           1000 non-null   int64
20  wifi                   1000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

```
In [5]: test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   battery_power    2000 non-null   int64
1   blue             2000 non-null   int64
2   clock_speed      2000 non-null   float64
3   dual_sim         2000 non-null   int64
4   fc               2000 non-null   int64
5   four_g           2000 non-null   int64
6   int_memory       2000 non-null   int64
7   m_dep            2000 non-null   float64
8   mobile_wt        2000 non-null   int64
9   n_cores          2000 non-null   int64
10  pc               2000 non-null   int64
11  px_height        2000 non-null   int64
12  px_width         2000 non-null   int64
13  ram              2000 non-null   int64
14  sc_h             2000 non-null   int64
15  sc_w             2000 non-null   int64
16  talk_time        2000 non-null   int64
17  three_g          2000 non-null   int64
18  touch_screen     2000 non-null   int64
19  wifi             2000 non-null   int64
20  price_range      2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
```

```
In [6]: x=train_df.drop('wifi',axis=1)
        y=train_df['wifi']
```

```
In [7]: x=test_df.drop('wifi',axis=1)
        y=test_df['wifi']
```

```
In [8]: train_df['dual_sim'].value_counts()
```

```
Out[8]: dual_sim
1      517
0      483
Name: count, dtype: int64
```

```
In [9]: test_df['blue'].value_counts()
```

```
Out[9]: blue
0      1010
1       990
Name: count, dtype: int64
```

```
In [11]: T={"Home Owner":{"Yes":1,"No":0}}
train_df=train_df.replace(T)
print(train_df)
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	
0	1	1043	1	1.8	1	14	0	5	\
1	2	841	1	0.5	1	4	1	61	
2	3	1807	1	2.8	0	1	0	27	
3	4	1546	0	0.5	1	18	1	25	
4	5	1434	0	1.4	0	11	1	49	
..	
995	996	1700	1	1.9	0	0	1	54	
996	997	609	0	1.8	1	0	0	13	
997	998	1185	0	1.4	0	1	1	8	
998	999	1533	1	0.5	1	0	0	50	
999	1000	1270	1	0.5	0	4	1	35	

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w	
0	0.1	193	...	16	226	1412	3476	12	7	\
1	0.8	191	...	12	746	857	3895	6	0	
2	0.9	186	...	4	1270	1366	2396	17	10	
3	0.5	96	...	20	295	1752	3893	10	0	
4	0.5	108	...	18	749	810	1773	15	8	
..	
995	0.5	170	...	17	644	913	2121	14	8	
996	0.9	186	...	2	1152	1632	1933	8	1	
997	0.5	80	...	12	477	825	1223	5	0	
998	0.4	171	...	12	38	832	2509	15	11	
999	0.1	140	...	19	457	608	2828	9	2	

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

```
In [12]: T={"Home Owner":{"Yes":1,"No":0}}
test_df=test_df.replace(T)
print(test_df)
```


	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	842	0	2.2	0	1	0	7 \
1	1021	1	0.5	1	0	1	53
2	563	1	0.5	1	2	1	41
3	615	1	2.5	0	0	0	10
4	1821	1	1.2	0	13	1	44
...
1995	794	1	0.5	1	0	1	2
1996	1965	1	2.6	1	0	0	39
1997	1911	0	0.9	1	1	1	36
1998	1512	0	0.9	0	4	1	46
1999	510	1	2.0	1	5	1	45

	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w
0	0.6	188	2	...	20	756	2549	9	7 \
1	0.7	136	3	...	905	1988	2631	17	3
2	0.9	145	5	...	1263	1716	2603	11	2
3	0.8	131	6	...	1216	1786	2769	16	8
4	0.6	141	2	...	1208	1212	1411	8	2
...
1995	0.8	106	6	...	1222	1890	668	13	4
1996	0.2	187	4	...	915	1965	2032	11	10
1997	0.7	108	8	...	868	1632	3057	9	1
1998	0.1	145	5	...	336	670	869	18	10
1999	0.9	168	6	...	483	754	3919	19	4

	talk_time	three_g	touch_screen	wifi	price_range
0	19	0	0	1	1
1	7	1	1	0	2
2	9	1	1	0	2
3	11	1	0	0	2
4	15	1	1	0	1
...
1995	19	1	1	0	0
1996	16	1	1	1	2
1997	5	1	1	0	3
1998	19	1	1	1	0
1999	2	1	1	1	3

[2000 rows x 21 columns]

```
In [13]: x=train_df.drop('wifi',axis=1)
y=train_df['wifi']
```

```
In [14]: x=test_df.drop('wifi',axis=1)
y=test_df['wifi']
```

```
In [15]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,train_size=0.7,random_state=42)
x_train.shape,x_test.shape
```

```
Out[15]: ((1400, 20), (600, 20))
```

```
In [16]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[16]: RandomForestClassifier()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In a Jupyter Environment ,please rerun this cell to show the HTML representation

```
In [17]: rf = RandomForestClassifier()
```

```
In [18]: params = {'max_depth': [2,3,5,10,20],
'min_samples_leaf': [5,10,20,50,100,200],
'n_estimators': [10,25,30,50,100,200]}
```

```
In [19]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rf,param_grid=params,cv = 2, scoring='accuracy')
grid_search.fit(x_train,y_train)
```

```
Out[19]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [2, 3, 5, 10, 20],
                                'min_samples_leaf': [5, 10, 20, 50, 100, 200],
                                'n_estimators': [10, 25, 30, 50, 100, 200]},
                    scoring='accuracy')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In a Jupyter Environment, please rerun this cell to show the HTML representation

```
In [20]: grid_search.best_score_
```

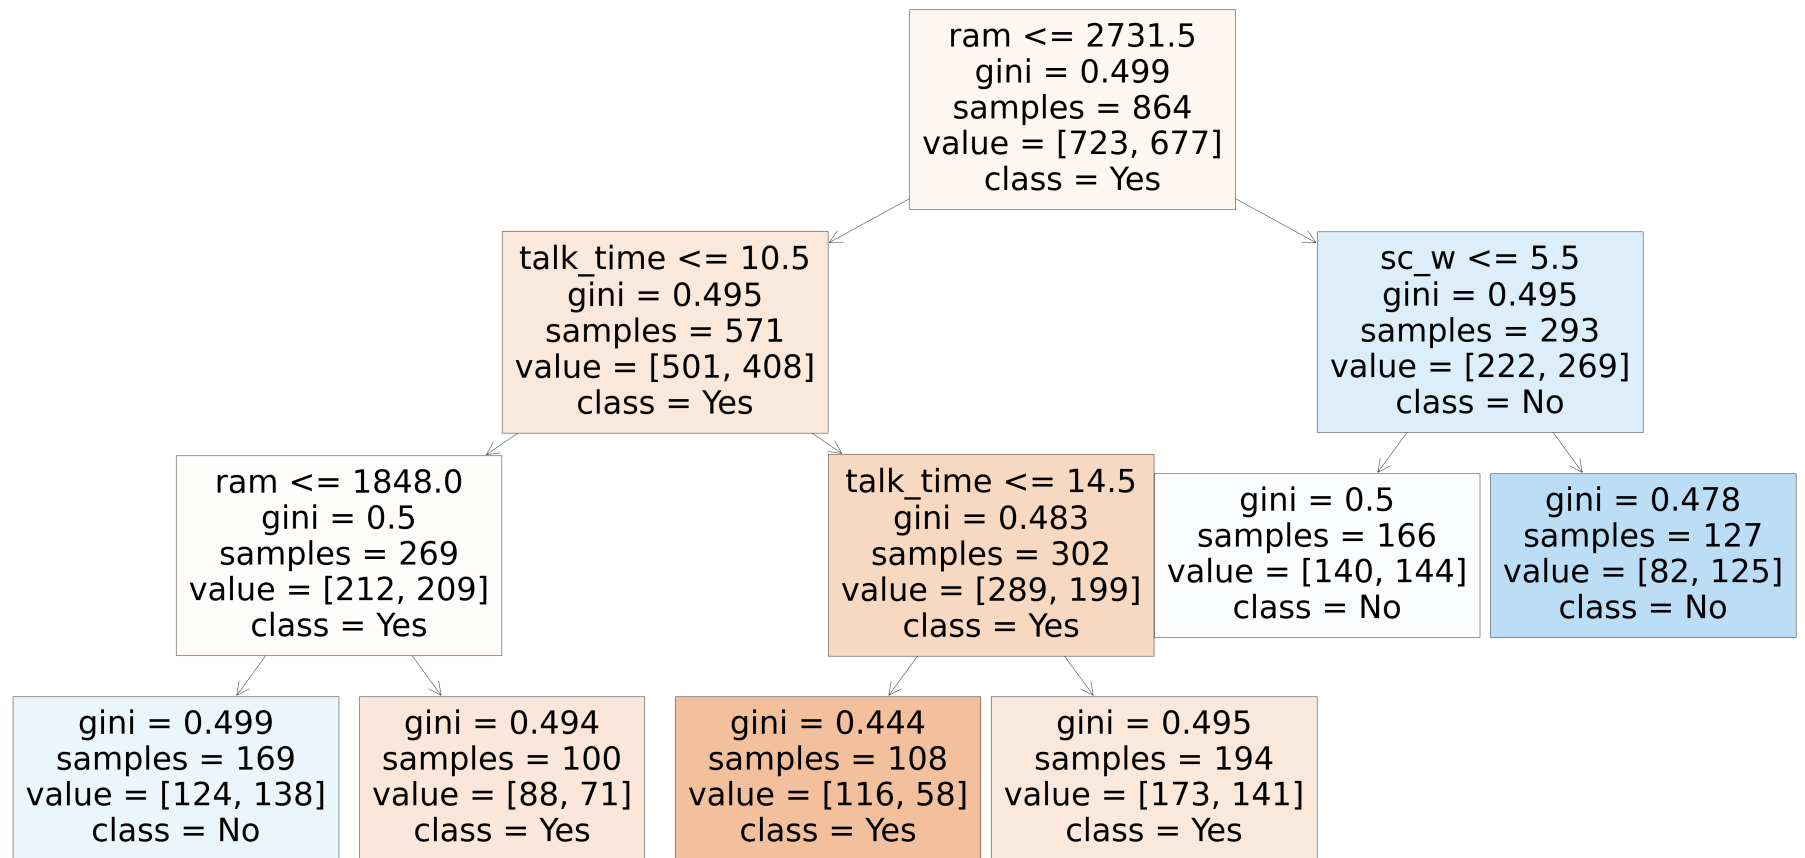
```
Out[20]: 0.5228571428571429
```

```
In [21]: rf_best = grid_search.best_estimator_
print(rf_best)
```

```
RandomForestClassifier(max_depth=5, min_samples_leaf=100, n_estimators=50)
```

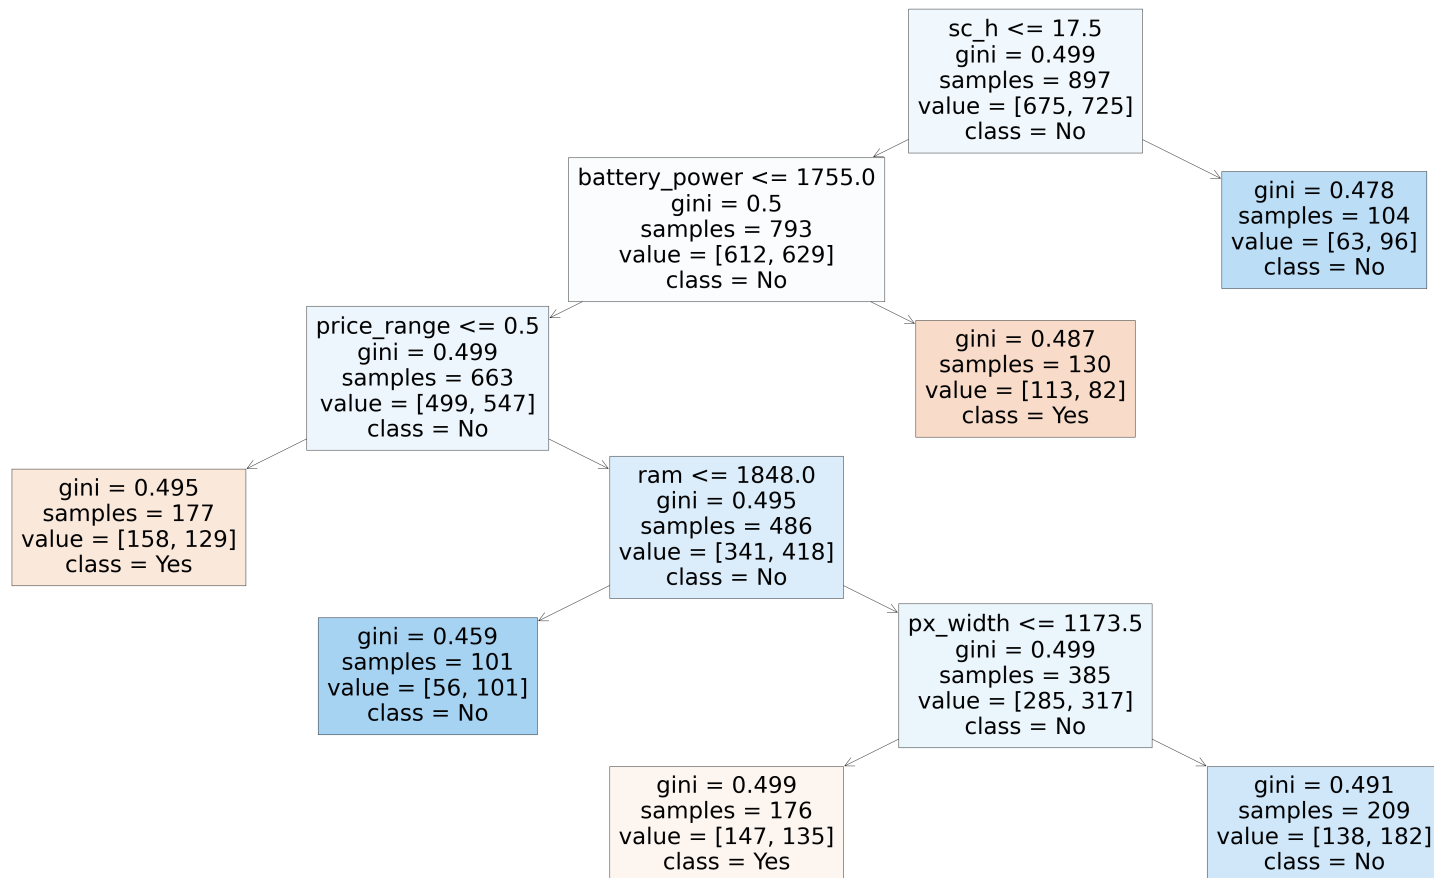
```
In [22]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[5], feature_names = x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[22]: [Text(0.5909090909090909, 0.875, 'ram <= 2731.5\ngini = 0.499\nsamples = 864\nvalue = [723, 677]\nclass = Yes'),
Text(0.36363636363636365, 0.625, 'talk_time <= 10.5\ngini = 0.495\nsamples = 571\nvalue = [501, 408]\nclass = Yes'),
Text(0.18181818181818182, 0.375, 'ram <= 1848.0\ngini = 0.5\nsamples = 269\nvalue = [212, 209]\nclass = Yes'),
Text(0.09090909090909091, 0.125, 'gini = 0.499\nsamples = 169\nvalue = [124, 138]\nclass = No'),
Text(0.2727272727272727, 0.125, 'gini = 0.494\nsamples = 100\nvalue = [88, 71]\nclass = Yes'),
Text(0.5454545454545454, 0.375, 'talk_time <= 14.5\ngini = 0.483\nsamples = 302\nvalue = [289, 199]\nclass = Yes'),
Text(0.45454545454545453, 0.125, 'gini = 0.444\nsamples = 108\nvalue = [116, 58]\nclass = Yes'),
Text(0.6363636363636364, 0.125, 'gini = 0.495\nsamples = 194\nvalue = [173, 141]\nclass = Yes'),
Text(0.8181818181818182, 0.625, 'sc_w <= 5.5\ngini = 0.495\nsamples = 293\nvalue = [222, 269]\nclass = No'),
Text(0.7272727272727273, 0.375, 'gini = 0.5\nsamples = 166\nvalue = [140, 144]\nclass = No'),
Text(0.9090909090909091, 0.375, 'gini = 0.478\nsamples = 127\nvalue = [82, 125]\nclass = No')]
```



```
In [23]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=["Yes","No"],filled=True)
```

```
Out[23]: [Text(0.6666666666666666, 0.9166666666666666, 'sc_h <= 17.5\ngini = 0.499\nsamples = 897\nvalue = [675, 725]\nclass = No'),
Text(0.5, 0.75, 'battery_power <= 1755.0\ngini = 0.5\nsamples = 793\nvalue = [612, 629]\nclass = No'),
Text(0.3333333333333333, 0.5833333333333334, 'price_range <= 0.5\ngini = 0.499\nsamples = 663\nvalue = [499, 547]\nclass = No'),
Text(0.16666666666666666, 0.41666666666666667, 'gini = 0.495\nsamples = 177\nvalue = [158, 129]\nclass = Yes'),
Text(0.5, 0.41666666666666667, 'ram <= 1848.0\ngini = 0.495\nsamples = 486\nvalue = [341, 418]\nclass = No'),
Text(0.3333333333333333, 0.25, 'gini = 0.459\nsamples = 101\nvalue = [56, 101]\nclass = No'),
Text(0.6666666666666666, 0.25, 'px_width <= 1173.5\ngini = 0.499\nsamples = 385\nvalue = [285, 317]\nclass = No'),
Text(0.5, 0.08333333333333333, 'gini = 0.499\nsamples = 176\nvalue = [147, 135]\nclass = Yes'),
Text(0.8333333333333334, 0.08333333333333333, 'gini = 0.491\nsamples = 209\nvalue = [138, 182]\nclass = No'),
Text(0.6666666666666666, 0.5833333333333334, 'gini = 0.487\nsamples = 130\nvalue = [113, 82]\nclass = Yes'),
Text(0.8333333333333334, 0.75, 'gini = 0.478\nsamples = 104\nvalue = [63, 96]\nclass = No')]
```



In [24]: rf_best.feature_importances_

Out[24]: array([0.07717307, 0.00879 , 0.04434522, 0.01953574, 0.03032118,
0.01928442, 0.08756478, 0.03099338, 0.05000981, 0.02300338,
0.06283458, 0.16018601, 0.07803434, 0.06519901, 0.05008368,
0.04774743, 0.10287966, 0.00509412, 0.00930888, 0.02761129])

```
In [25]: imp_df = pd.DataFrame({"Vername": x_train.columns, "Imp": rf_best.feature_importances_})
imp_df.sort_values(by="Imp", ascending=False)
```

Out[25]:

	Vername	Imp
11	px_height	0.160186
16	talk_time	0.102880
6	int_memory	0.087565
12	px_width	0.078034
0	battery_power	0.077173
13	ram	0.065199
10	pc	0.062835
14	sc_h	0.050084
8	mobile_wt	0.050010
15	sc_w	0.047747
2	clock_speed	0.044345
7	m_dep	0.030993
4	fc	0.030321
19	price_range	0.027611
9	n_cores	0.023003
3	dual_sim	0.019536
5	four_g	0.019284
18	touch_screen	0.009309
1	blue	0.008790
17	three_g	0.005094

